



ML Project

Face Mask Detection

Group 3

Team Member: Huishan Zhang, Linjing Li, Xi Fei, Zhujun Tian

AGENDA

Part 1 Project Introduction

Part 2 Dataset Characteristics

Part 3 Convolutional Neural Network

Part 4 Data Processing and Training

Part 5 Testing and Evaluation

Part 6 Conclusion



Part 1 Project Introduction

Project introduction



Project introduction



- Face Mask-wearing is an efficient way to help control Covid-19 outbreak
- As Face Mask is became essential for everyone while roaming outside, identifying mask wear help monitor public behavior and contribute towards constraining the COVID-19 pandemic.

Part 2 Dataset Characteristics

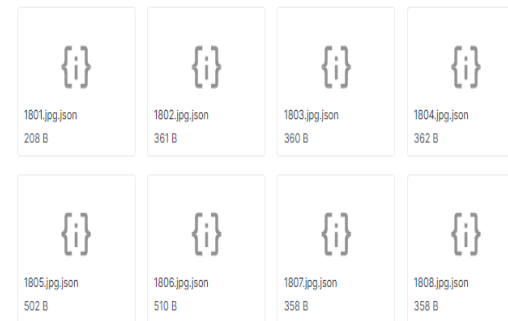
Dataset Characteristics

Data Explorer

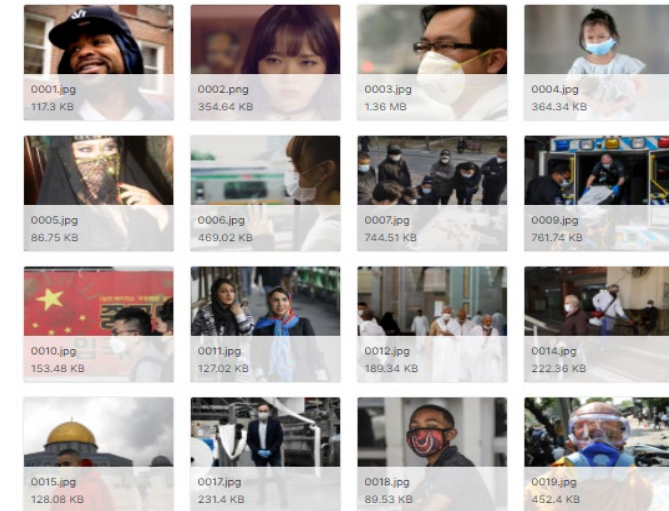
2.58 GB

- Medical mask
 - Medical mask
 - Medical Mask
 - annotations
 - images
 - meta.json
- submission.csv
- train.csv

< annotations (4326 files)



< images (6024 files)



< train.csv (571.71 KB)



Detail Compact Column

6 of 6 columns

# name	# x1	# x2	# y1	# y2	# classname
2756 . png	69	126	294	392	face_with_mask

Dataset Characteristics

- **train.csv**: take the first record from `meta_data(train.csv)`; then extract all the records with the same name ('2756.png') from `meta_data(train.csv)`
- **images file**: plot the images with the same name ('2756.png')
- **annotations file**: detail information regarding FileName '2756.png'

```
meta_name=meta_data["name"][0]
meta_name
```

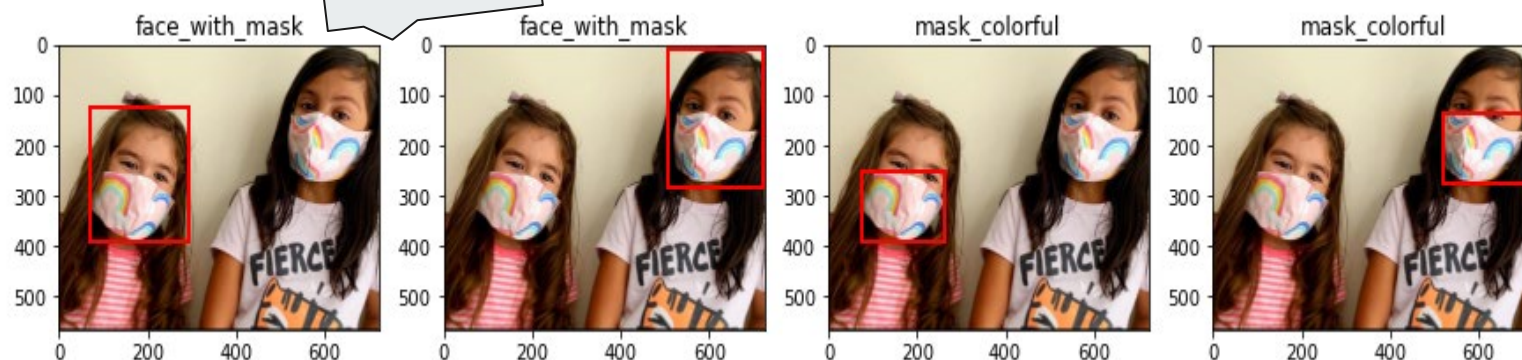
```
'2756.png'
```

```
meta_data[meta_data["name"]==meta_name]
```

	name	x1	x2	y1	y2	classname
0	2756.png	69	126	294	392	face_with_mask
1	2756.png	505	10	723	283	face_with_mask
2	2756.png	75	252	264	390	mask_colorful
3	2756.png	521	136	711	277	mask_colorful

```
jsonfiles = []
for i in os.listdir(directory):
    jsonfiles.append(getJSON(os.path.join(directory,i)))
jsonfiles[878]
```

```
{'FileName': '2756.png',
 'NumOfAnno': 4,
 'Annotations': [{'isProtected': False,
 'ID': 598039385457921920,
 'BoundingBox': [69, 126, 294, 392],
 'classname': 'face_with_mask',
 'Confidence': 1,
 'Attributes': {}},
 {'isProtected': False,
 'ID': 702410169516003712,
 'BoundingBox': [505, 10, 723, 283],
 'classname': 'face_with_mask',
 'Confidence': 1,
 'Attributes': {}},
 {'isProtected': False,
 'ID': 152706552218535680,
 'BoundingBox': [75, 252, 264, 390],
 'classname': 'mask_colorful',
 'Confidence': 1,
 'Attributes': {}},
 {'isProtected': False,
 'ID': 3262639018530855,
 'BoundingBox': [521, 136, 711, 277],
 'classname': 'mask_colorful',
 'Confidence': 1,
 'Attributes': {}}]}
```




Dataset Characteristics - train.csv(meta_data)

```
meta_data.head()
```

	name	x1	x2	y1	y2	classname
0	2756.png	69	126	294	392	face_with_mask
1	2756.png	505	10	723	283	face_with_mask
2	2756.png	75	252	264	390	mask_colorful
3	2756.png	521	136	711	277	mask_colorful
4	6098.jpg	360	85	728	653	face_no_mask

```
print(len(meta_data))
```

15412



```
## Counts of face_with_mask and face_no_mask
```

```
face_with_mask=train_data[train_data["classname"]=="face_with_mask"]
```

```
face_no_mask=train_data[train_data["classname"]=="face_no_mask"]
```

```
print("count of face with mask: "+str(len(face_with_mask))+"\ncount of face no mask: "+str(len(face_no_mask)))
```

count of face with mask: 4180

count of face no mask: 1569



Part 3

Convolutional Neural Network

Convolutional Neural Network(CNN)

Technology that teaches computer to 'see'

Application: Image Recognition, Autonomous Driving, Teach Computer to Play Video Games/ Cooking...




157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

What the computer sees

157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

An image is just a matrix of numbers [0,255]!
i.e., 1080x1080x3 for an RGB image



Convolutional Neural Network(CNN)

What make the recognition so hard for the computer?



Viewpoint variation




Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter



Intra-class variation



Solution?

Feature Detector/Filter/Kernel

CNN Key Concepts and Architecture -- Filtering

Filters are common concept in the image processing field. They are usually number matrix that can detect horizontal, vertical, edges... of images. But in CNNs, filters are not defined. The value of each filter is learned during the training process.

By being able to learn the values of different filters, CNNs can find more meaning from images that humans and human designed filters might not be able to find.

-1	-2	-1
0	0	0
1	2	1


Horizontal

-1	0	1
-2	0	2
-1	0	1

Vertical



Low level features



Edges, dark spots

Mid level features



Eyes, ears, nose

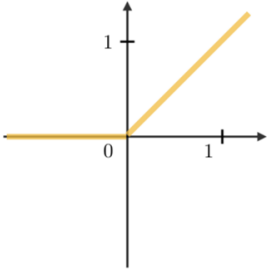
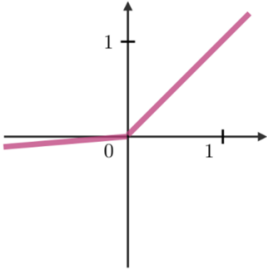
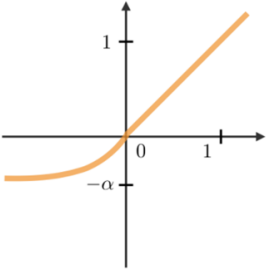
High level features



Facial structure

CNN Key Concepts and Architecture-- Commonly Used Activation Functions

□ **Rectified Linear Unit** — The rectified linear unit layer (ReLU) is an activation function g that is used on all elements of the volume. It aims at introducing non-linearities to the network. Its variants are summarized in the table below:

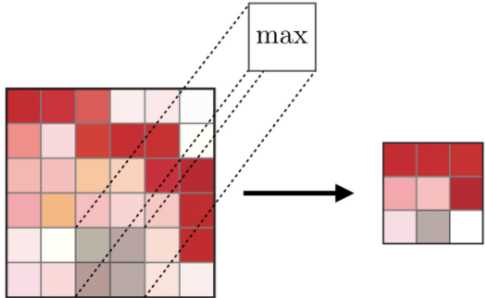
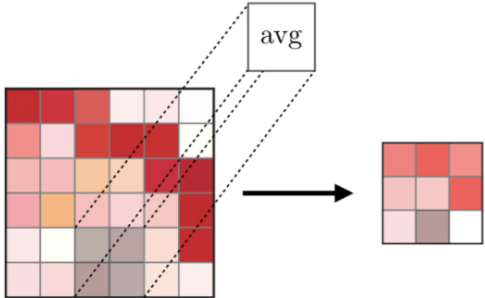
ReLU	Leaky ReLU	ELU
$g(z) = \max(0, z)$	$g(z) = \max(\epsilon z, z)$ with $\epsilon \ll 1$	$g(z) = \max(\alpha(e^z - 1), z)$ with $\alpha \ll 1$
		
<ul style="list-style-type: none">• Non-linearity complexities biologically interpretable	<ul style="list-style-type: none">• Addresses dying ReLU issue for negative values	<ul style="list-style-type: none">• Differentiable everywhere



□ **Softmax** — The softmax step can be seen as a generalized logistic function that takes as input a vector of scores $x \in \mathbb{R}^n$ and outputs a vector of output probability $p \in \mathbb{R}^n$ through a softmax function at the end of the architecture. It is defined as follows:

$$p = \begin{pmatrix} p_1 \\ \vdots \\ p_n \end{pmatrix} \quad \text{where} \quad p_i = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

CNN Key Concepts and Architecture -- Pooling


The pooling layer is a downsampling operation, typically applied after a convolution layer, which does some spatial invariance. In particular, max and average pooling are special kinds of pooling where the maximum and average value is taken, respectively.

Type	Max pooling	Average pooling
Purpose	Each pooling operation selects the maximum value of the current view	Each pooling operation averages the values of the current view
Illustration		
Comments	<ul style="list-style-type: none">• Preserves detected features• Most commonly used	<ul style="list-style-type: none">• Downsamples feature map• Used in LeNet





CNN Key Concepts and Architecture -- Fully Connected

Every value gets a vote




Convolutional Neural Network Architecture



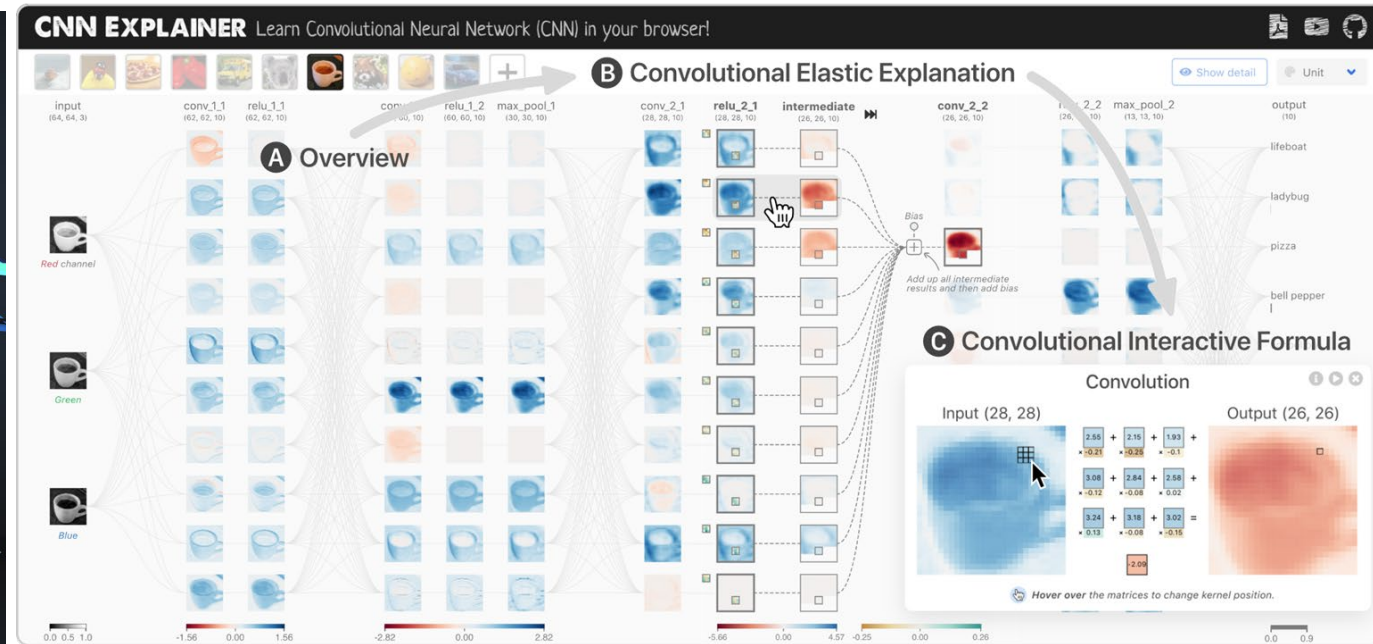
Visualization of Convolutional Neural Networks

(Adam Harley & Zijie Wang)

<https://www.cs.ryerson.ca/~aharley/vis/conv/>



<https://zijie.wang/papers/cnn-explainer/>



Convolutional Neural Network(CNN)


Concept to code

```
import tensorflow as tf

def generate_model():
    model = tf.keras.Sequential([
        # first convolutional layer
        tf.keras.layers.Conv2D(32, filter_size=3, activation='relu'),
        tf.keras.layers.MaxPool2D(pool_size=2, strides=2),

        # second convolutional layer
        tf.keras.layers.Conv2D(64, filter_size=3, activation='relu'),
        tf.keras.layers.MaxPool2D(pool_size=2, strides=2),

        # fully connected classifier
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(1024, activation='relu'),
        tf.keras.layers.Dense(10, activation='softmax') # 10 outputs
    ])
    return model
```



Part 4 Data Processing and Training

Data Processing

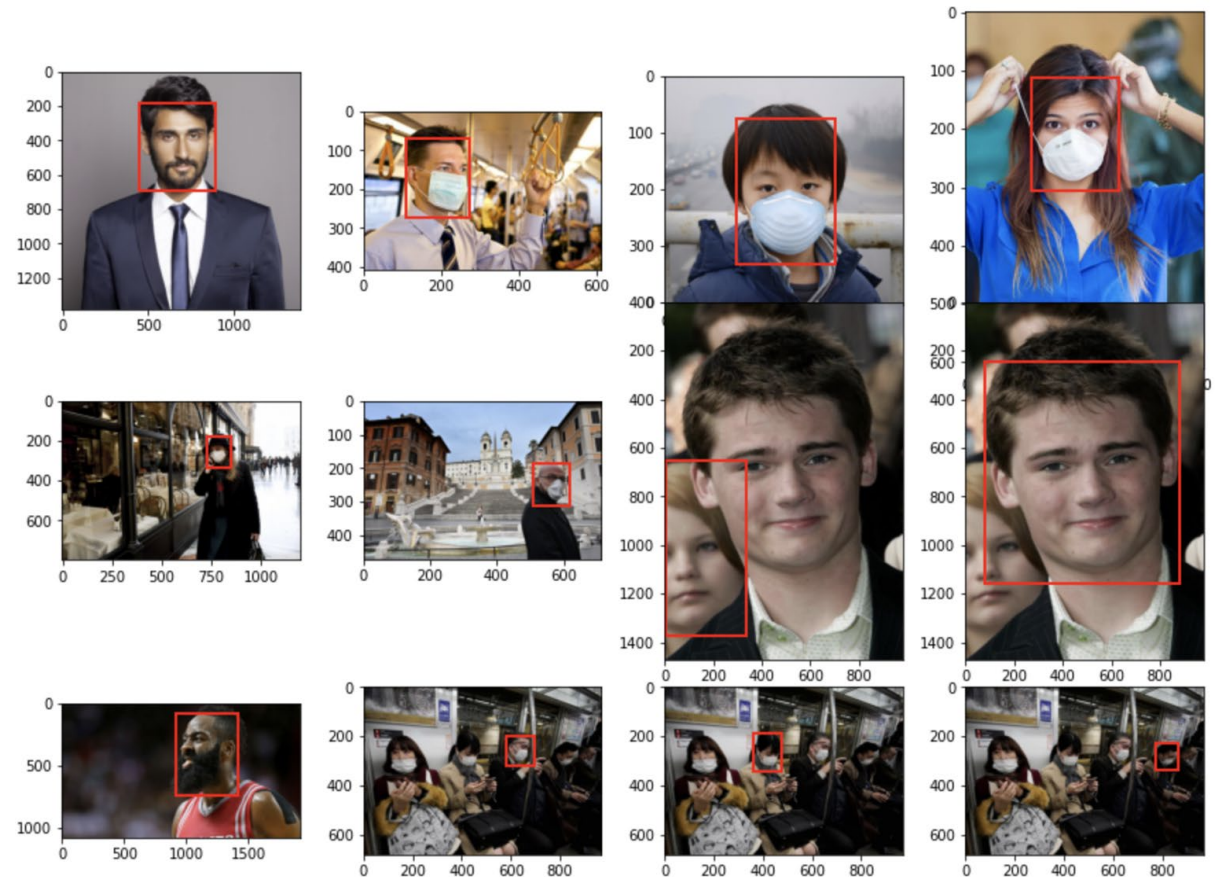
```
df = meta_data[(meta_data.classname=='face_with_mask') | (meta_data.classname=='face_no_mask')]
```

df.head()

	name	x1	x2	y1	y2	classname
0	1801.jpg	451	186	895	697	face_no_mask
1	1802.jpg	110	71	273	272	face_with_mask
2	1803.jpg	126	75	303	333	face_with_mask
3	1804.jpg	112	113	262	307	face_with_mask
4	1805.jpg	728	180	853	336	face_with_mask

Total images: 3390

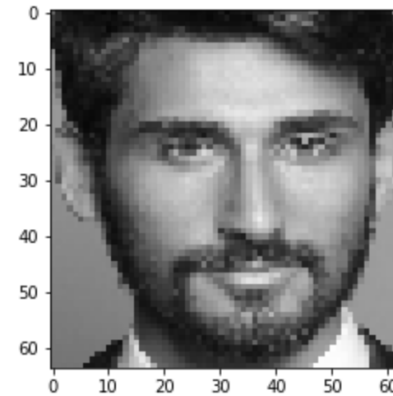
Total faces: 5749



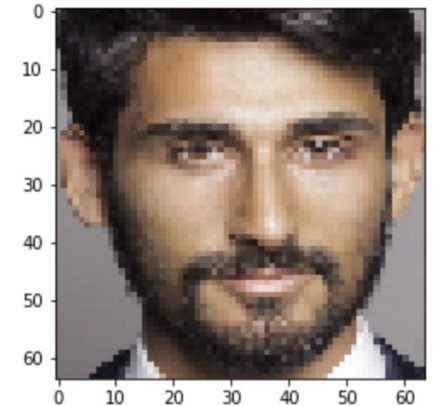
Data Processing

```
[17] def img_to_cv2(path, x1, x2, y1, y2, img_rows, img_cols, color_type=1):  
    img_arr = []  
    # Load as grayscale  
    if color_type == 1:  
        img_arr = cv2.imread(path, 0)  
    elif color_type == 3:  
        img_arr = cv2.imread(path)  
  
    # extract face image  
    img_arr = img_arr[x2:y2, x1:y1]  
    # Reshape size  
    resized = cv2.resize(img_arr, (img_cols, img_rows))  
  
    return resized
```

Grayscale (64, 64, 1)



RGB(64, 64, 3)



1. Read Image

- The output of `cv2.imread()` is an array of BGR(Blue, Green, Red) values. (Use `cv2.cvtColor` to convert BGR to RGB for plotting)
- Use grayscale as an input.

2. Extract face

3. Reshape face as input size

```
▶ x[0]  
[ ] array([[100,  41,  20, ..., 139, 160, 160],  
          [ 59,  44,  25, ...,  40, 160, 159],  
          [ 63,  19,  28, ...,  29, 140, 168],  
          ...,  
          [118, 117, 118, ...,  35, 144, 150],  
          [117, 117, 118, ...,  29,  45,  39],  
          [116, 116, 115, ...,  28,  33,  48]], dtype=uint8)
```

Build model

Three convolution layers

```
def create_cnn_model(img_rows, img_cols, color_type=1):
    model = Sequential()
    model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(img_rows, img_cols, color_type)))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.2))

    model.add(Conv2D(64, (3, 3), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.2))

    model.add(Conv2D(128, (3, 3), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.2))

    model.add(Flatten())
    model.add(Dense(2, activation='softmax'))

    model.summary()

    return model
```

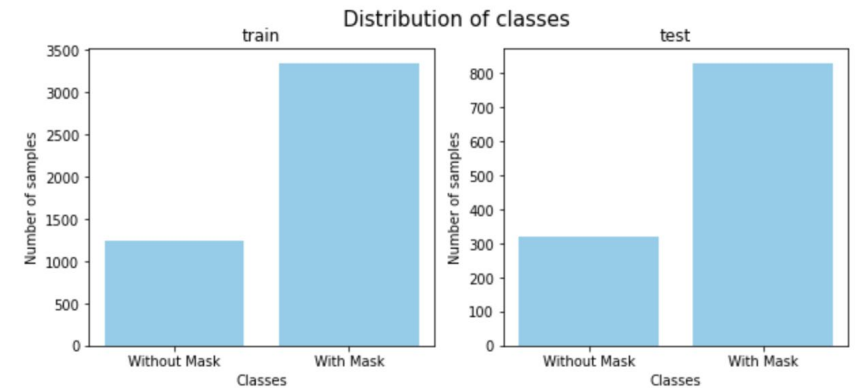
Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 62, 62, 32)	320
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
dropout (Dropout)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 29, 29, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 64)	0
dropout_1 (Dropout)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 12, 12, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
dropout_2 (Dropout)	(None, 6, 6, 128)	0
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 2)	9218
=====		
Total params: 101,890		
Trainable params: 101,890		
Non-trainable params: 0		

Train model

- Normalize & split dataset

```
def normalize_data(X, y, img_rows, img_cols, color_type=1):  
    X = np.array(X).reshape(-1, img_rows, img_cols, color_type)  
    y = np.array(y)  
    y = tf.keras.utils.to_categorical(y)  
    X = X.astype('float32') / 255  
  
    return X, y
```



- Set checkpoint to save the best model

```
[ ] def train_model(model, X_train, y_train, epochs, batch_size, checkpoint=None):  
    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])  
    history = model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size, validation_split=0.2, shuffle=True, callbacks=[checkpoint])  
    return history
```

```
[ ] checkpoint_path = '/checkpoint'  
    checkpoint = ModelCheckpoint(checkpoint_path, monitor='val_loss', verbose=0, save_best_only=True, mode='auto')
```

Train model

```
[5] # parameter
    img_rows, img_cols = 64, 64
    color_type = 1
    random_state = 45
    epochs = 20
    batch_size = 32
```

- train model

```
▶ model = create_cnn_model(img_rows, img_cols, color_type)
  history = train_model(model, X_train, y_train, epochs, batch_size, checkpoint)
```

- load checkpoint to model


```
▶ model.load_weights(checkpoint_path)
```

```
Epoch 1/20
115/115 [=====] - 46s 247ms/step - loss: 0.5732 - accuracy: 0.7223 - val_loss: 0.5413 - val_accuracy: 0.7217
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 2/20
115/115 [=====] - 28s 240ms/step - loss: 0.4921 - accuracy: 0.7390 - val_loss: 0.4128 - val_accuracy: 0.7750
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 3/20
115/115 [=====] - 28s 241ms/step - loss: 0.4286 - accuracy: 0.7835 - val_loss: 0.3670 - val_accuracy: 0.8489
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 4/20
115/115 [=====] - 28s 241ms/step - loss: 0.3607 - accuracy: 0.8298 - val_loss: 0.3238 - val_accuracy: 0.8696
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 5/20
115/115 [=====] - 28s 245ms/step - loss: 0.3324 - accuracy: 0.8527 - val_loss: 0.3269 - val_accuracy: 0.8772
Epoch 6/20
115/115 [=====] - 28s 246ms/step - loss: 0.3216 - accuracy: 0.8517 - val_loss: 0.2691 - val_accuracy: 0.8957
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 7/20
115/115 [=====] - 28s 246ms/step - loss: 0.2778 - accuracy: 0.8813 - val_loss: 0.2575 - val_accuracy: 0.8957
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 8/20
115/115 [=====] - 28s 243ms/step - loss: 0.2376 - accuracy: 0.9033 - val_loss: 0.2238 - val_accuracy: 0.9022
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 9/20
115/115 [=====] - 28s 244ms/step - loss: 0.2219 - accuracy: 0.8987 - val_loss: 0.2640 - val_accuracy: 0.8783
Epoch 10/20
115/115 [=====] - 28s 243ms/step - loss: 0.2170 - accuracy: 0.9100 - val_loss: 0.2690 - val_accuracy: 0.8674
Epoch 11/20
115/115 [=====] - 28s 242ms/step - loss: 0.2030 - accuracy: 0.9116 - val_loss: 0.2153 - val_accuracy: 0.9196
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 12/20
115/115 [=====] - 28s 241ms/step - loss: 0.1713 - accuracy: 0.9292 - val_loss: 0.1967 - val_accuracy: 0.9261
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 13/20
115/115 [=====] - 28s 243ms/step - loss: 0.1574 - accuracy: 0.9311 - val_loss: 0.2198 - val_accuracy: 0.8967
Epoch 14/20
115/115 [=====] - 28s 242ms/step - loss: 0.1616 - accuracy: 0.9318 - val_loss: 0.1998 - val_accuracy: 0.9152
Epoch 15/20
115/115 [=====] - 28s 243ms/step - loss: 0.1510 - accuracy: 0.9432 - val_loss: 0.1897 - val_accuracy: 0.9250
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 16/20
115/115 [=====] - 28s 243ms/step - loss: 0.1222 - accuracy: 0.9517 - val_loss: 0.2162 - val_accuracy: 0.9011
Epoch 17/20
115/115 [=====] - 28s 243ms/step - loss: 0.1245 - accuracy: 0.9537 - val_loss: 0.1795 - val_accuracy: 0.9207
INFO:tensorflow:Assets written to: /checkpoint/assets
Epoch 18/20
115/115 [=====] - 28s 243ms/step - loss: 0.1149 - accuracy: 0.9578 - val_loss: 0.2347 - val_accuracy: 0.9022
Epoch 19/20
115/115 [=====] - 28s 245ms/step - loss: 0.1225 - accuracy: 0.9507 - val_loss: 0.2016 - val_accuracy: 0.9228
Epoch 20/20
115/115 [=====] - 28s 244ms/step - loss: 0.1152 - accuracy: 0.9536 - val_loss: 0.2093 - val_accuracy: 0.9163
```



Part 5 Testing and Evaluation

Visualize training and evaluate testing set

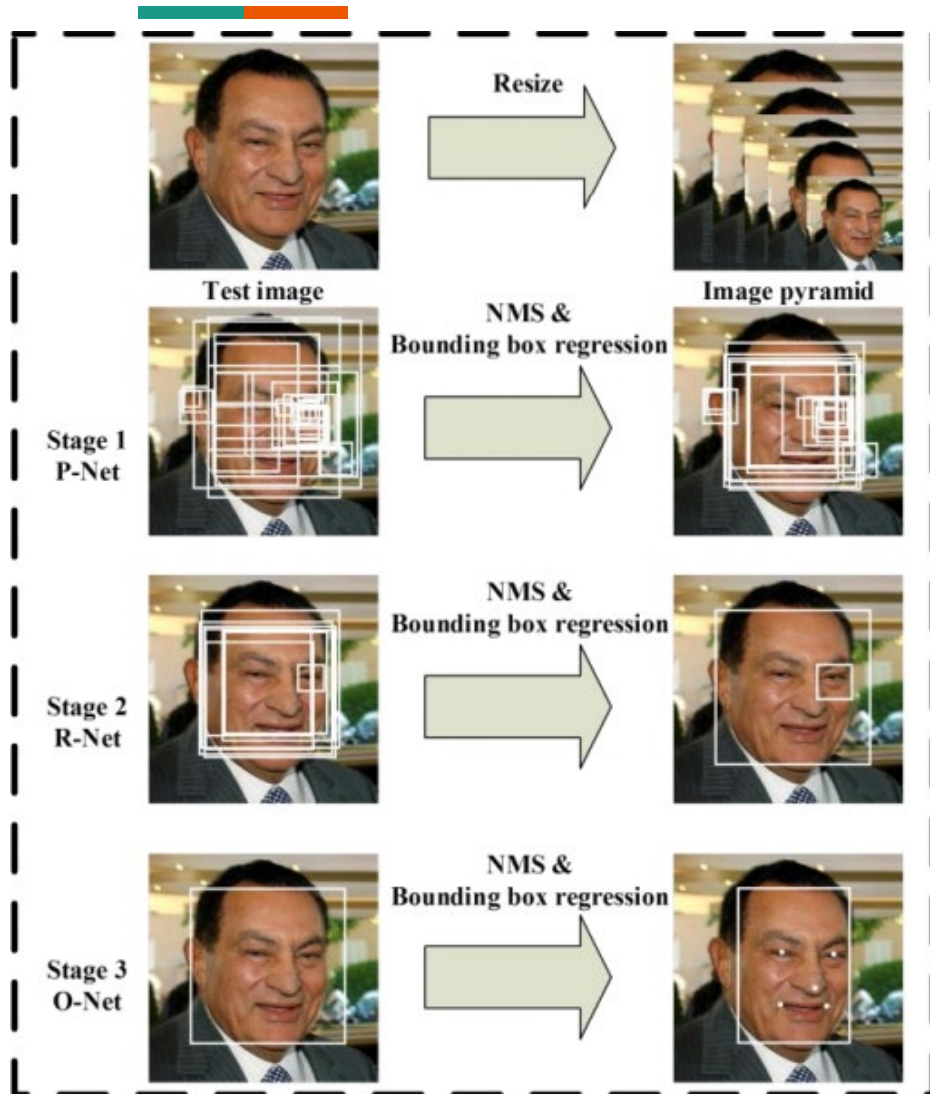


```
train_score = model.evaluate(X_train, y_train, batch_size=batch_size)
test_score = model.evaluate(X_test, y_test, batch_size=batch_size)
```



```
144/144 [=====] - 9s 59ms/step - loss: 0.1024 - accuracy: 0.9630
36/36 [=====] - 2s 59ms/step - loss: 0.1941 - accuracy: 0.9252
```

Single image prediction - MTCNN

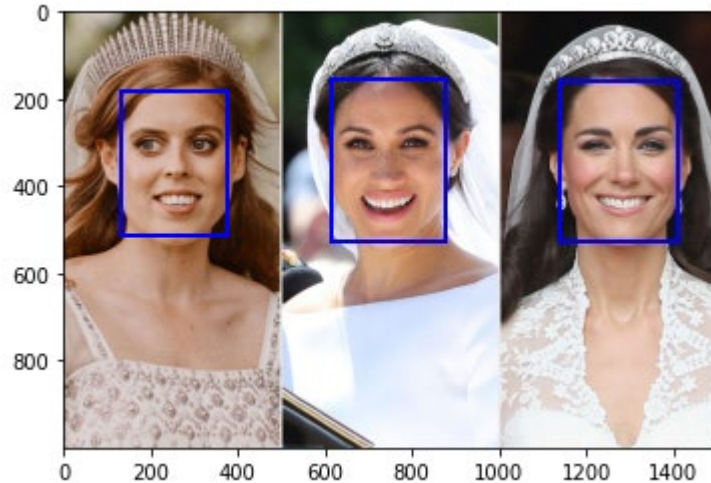


First Stage: P-Net(Proposal Net). Create multiple frames scan through the whole image.

Second Stage: R-Net(Refinement Net). Use information from P-Net as input, frames which do not contain faces will be rejected by the R-Net.

Third Stage: O-Net(Output Net). O-Net will eventually outputs the face and facial landmarks position detecting from the image.

Single image prediction



```
def face_detect(path):  
    box_arr = []  
    img = plt.imread(path)  
    faces = MTCNN().detect_faces(img)  
    for face in faces:  
        box_arr.append(face['box'])  
    return box_arr
```


```
file_name = "test.jpeg"  
path = file_name  
  
fig, ax = plt.subplots()  
img = plt.imread(path)  
ax.imshow(img)  
  
# x1, y1, width, height = face['box']  
faces = face_detect(path)  
for i, face in enumerate(faces):  
    x1, y1, width, height = face[0], face[1], face[2], face[3]  
    face = patches.Rectangle((x1, y1), width, height, linewidth=2, edgecolor='b', facecolor='none')  
    ax.add_patch(face)  
  
plt.show()
```

```
data_1 = load_test_data(path, img_rows, img_cols, color_type)  
data_1 = normalize_test_data(data_1, img_rows, img_cols, color_type)
```

```
model.predict(data_1)
```

```
array([[9.8476613e-01, 1.5233894e-02],  
       [9.9991572e-01, 8.4258529e-05],  
       [9.9990559e-01, 9.4352967e-05]], dtype=float32)
```

Single image prediction



```
file_name = "test2.jpeg"
path = file_name

fig, ax = plt.subplots()
img = plt.imread(path)
ax.imshow(img)

# x1, y1, width, height = face['box']
faces = face_detect(path)
for i, face in enumerate(faces):
    x1, y1, width, height = face[0], face[1], face[2], face[3]
    face = patches.Rectangle((x1, y1), width, height, linewidth=2, edgecolor='b', facecolor='none')
    ax.add_patch(face)

plt.show()
```

```
data_2 = load_test_data(path, img_rows, img_cols, color_type)
data_2 = normalize_test_data(data_2, img_rows, img_cols, color_type)
```

```
model.predict(data_2)
```

```
array([[0.01019729, 0.9898028 ]], dtype=float32)
```



Part 6 Conclusion



Thank you

References



How Convolutional Neural Networks Work (CNNs Explained & Visualized)

<https://www.youtube.com/watch?v=pj9-rr1wDhM>

How Convolutional Neural Networks work

<https://www.youtube.com/watch?v=FmpDIaiMleA>

MIT 6.S191 (2020): Convolutional Neural Networks

<https://www.youtube.com/watch?v=iaSUYvmCekI>

Convolutional Neural Networks cheatsheet

<https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-convolutional-neural-networks>

An Interactive Node-Link Visualization of Convolutional Neural Networks

<https://www.cs.ryerson.ca/~aharley/vis/conv/>

Different Kinds of Convolutional Filters

<https://www.saama.com/different-kinds-convolutional-filters/>

How convolutional neural networks work, in depth

https://www.youtube.com/watch?v=JB8T_zN7ZC0

CNN Explainer: Learning Convolutional Neural Networks with Interactive Visualization

<https://zijie.wang/papers/cnn-explainer/>

<https://www.kaggle.com/dabawse/detecting-face-masks-with-5-models/data>